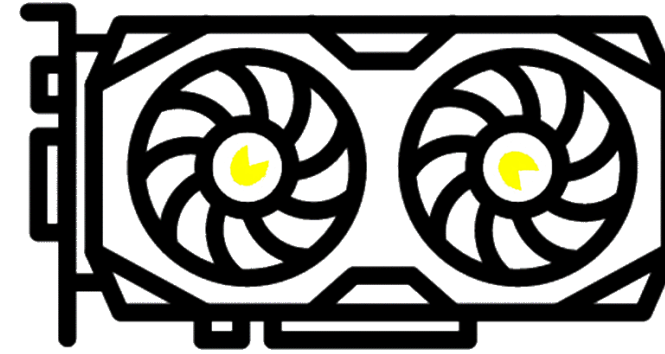


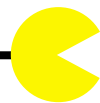
Parallel Computing

GPU Program Flow (like a boss)



CUDA Streams

- A sequence of operations executed on the device **in order** as executed on the host
- The operations (kernels and data transfers) in a stream **cannot** overlap
- The default stream:
 - Synchronizing stream with respect to operations on the device on **any other stream**
 - Operation starts when all previously issued operations in any stream are finished
 - New launched operations begin after the default stream operation is finished
- Non-default stream:
 - All operations are async (non-blocking)



CUDA Streams – How to

- Create a (or N) stream(s)


```
1  cudaStream_t stream;  
2  gpuErrCheck(cudaStreamCreate(&stream));           // gpuErrCheck is our PAC macro
```

- Execute a kernel

```
3  packKernel<<< blocks, threads, SHMbytes >>>();           // default stream  
4  packKernel<<< blocks, threads, SHMbytes, stream >>>();    // use a non-default stream  
5  gpuErrCheck(cudaPeekAtLastError());                // did the kernel launch work?
```

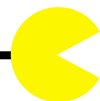
- We are not mad and cleanup our mess

```
6  gpuErrCheck(cudaStreamDestroy(stream));
```




CUDA Streams

- Since everything is executed async in respect to the host, we need to synchronize
- `cudaDeviceSynchronize()` wait until all previous issued operations **in all streams** are done
- `cudaStreamSynchronize(stream)`
wait until all previous issued operations **in this stream** are done
- `cudaStreamQuery(stream)` check if all operations in this stream are finished (empty steam)
- `cudaEventSynchronize(event)` and `cudaEventQuery(event)` - see code of last week
- `cudaStreamWaitEvent(event)`
can sync on a specific event of any stream, even of another device




CUDA Streams – Some notes

- CUDA 7 has a major improvement: `--default-stream per-thread` compile argument
- Every thread gets its own default stream which **is a non-default-stream**
 - No global device sync
- Handy for OMP parallel directives as no tracking of streams needs to be done
- You must implement a domain decomposition of data and processing
- Don't overdo it! A few streams are most often more than enough!



Observations

- Using multiple streams can increase the occupancy of the hardware significantly
- The streams need different CPU-threads in order to run in parallel
- There are free resources on the CPU:
Maybe one of the tasks can be done on the CPU instead of the GPU?



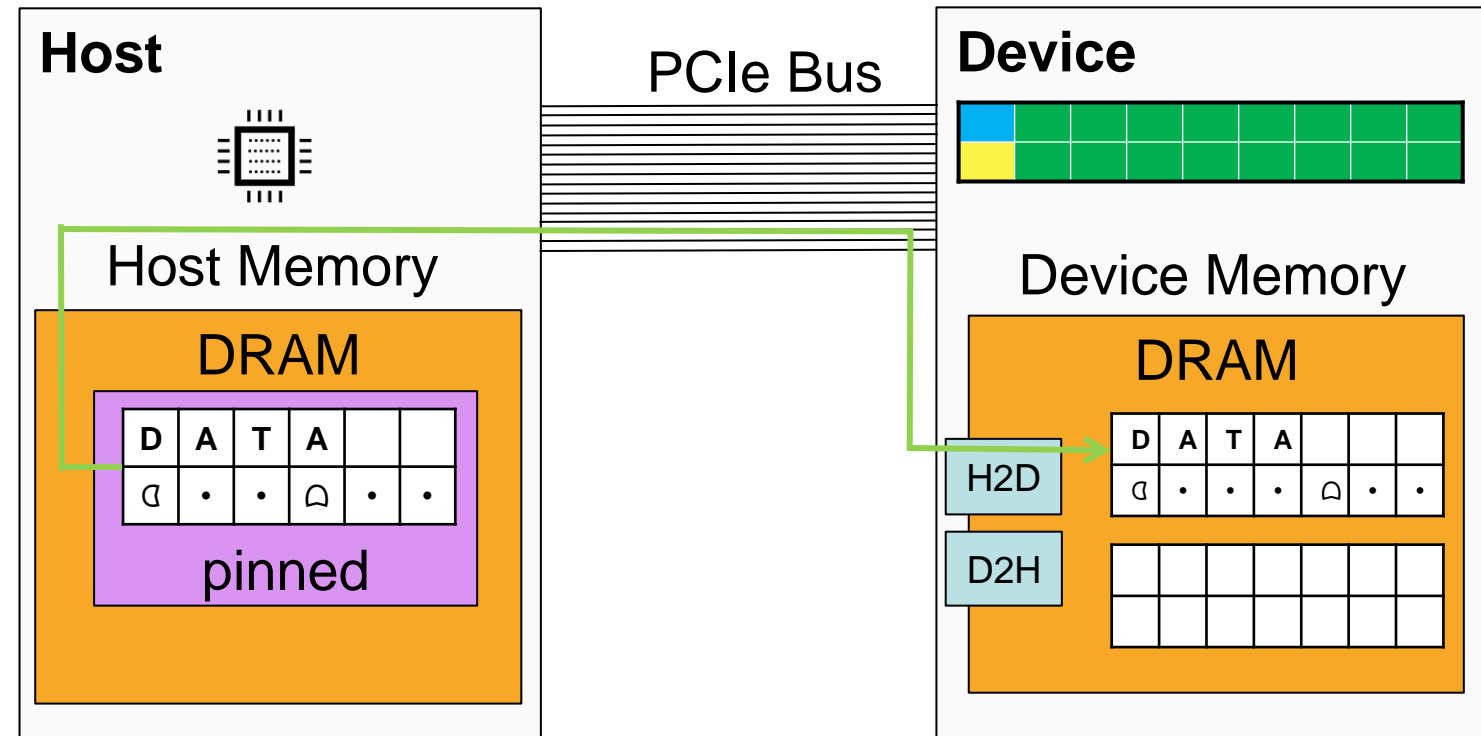
CUDA Processing Flow using async data transfer

1. Load data into Host Memory
 - CPU load
 - Needs to be pinned memory

```
int* h_matrixA;
cudaMallocHost(&h_matrixA,
               size * sizeof(int));
```

2. Copy data to Device using H2D (async)
 - H2D engine load

```
cudaMemcpyAsync(d_matrixA,
                h_matrixA,
                size * sizeof(int),
                cudaMemcpyHostToDevice,
                stream);
```

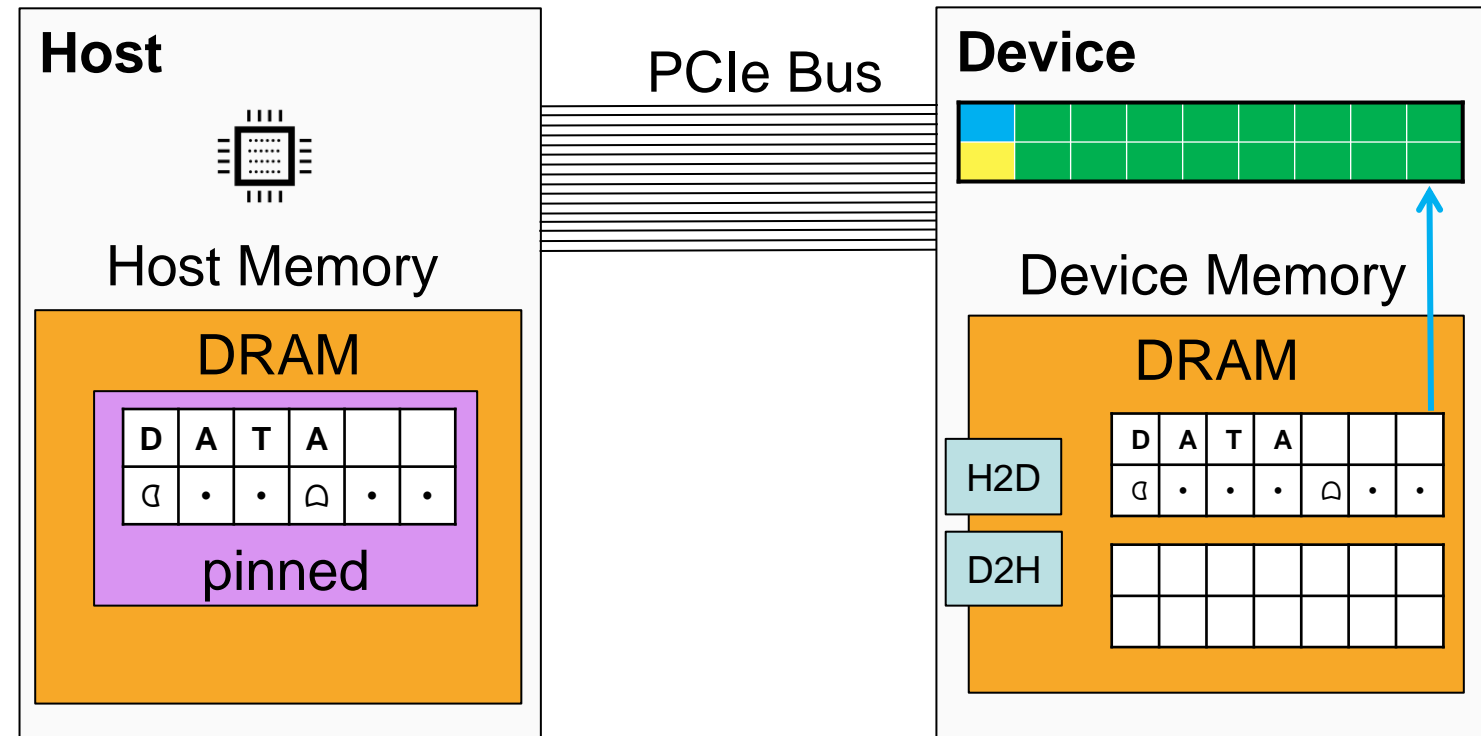


CUDA Processing Flow using async data transfer

3. Execute kernel
 - GPU load (kernel engine)


```
packKernel<<< blocks,
               threads,
               SHMbytes,
               stream >>>(...);
```
4. Same way back using async D2H ☺
5. Clean up your mess


```
cudaFree(d_matrixA);
cudaFreeHost(h_matrixA);
...
```


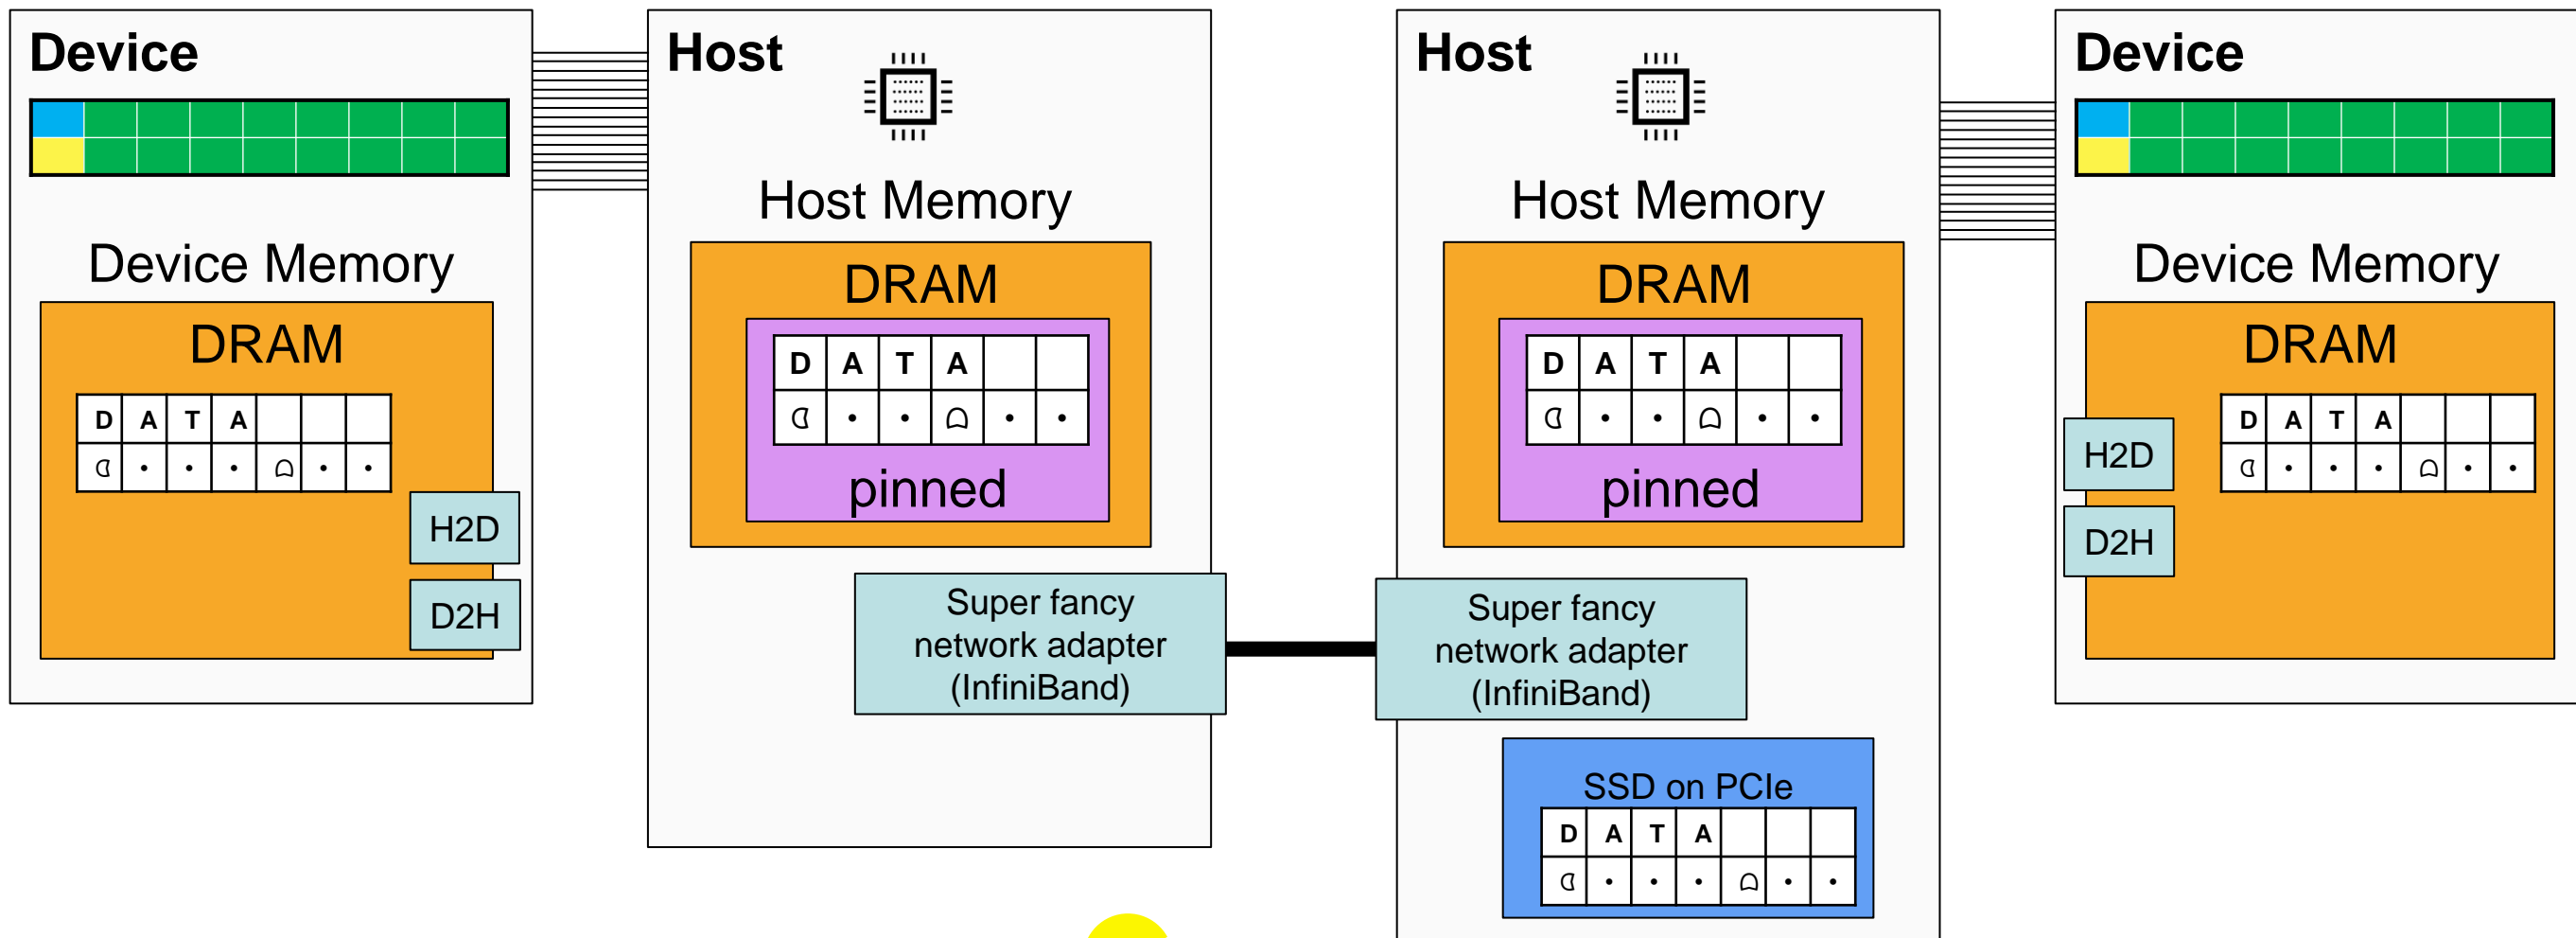


Observations

- Offload CPU work to DMA engines by using async copy (and thus pinned memory)
- Think hard how to use these free CPU cycles in parallel and efficient
Maybe do some fancy AVX stuff 😊




CUDA Processing Flow – think outside of the box - RDMA




CUDA Streams – Some more notes

- In real life cases, the kernel uses more time than the copy
 - Even one CPU thread using multiple streams can overlap H2D/D2H with the kernels
 - You can prepare work in advance
 - Use CUDA events to synchronize/wait at the right spots
- Use the free CPU cores to:
 - Proceed with the GPU results (preferred)
 - Do the same thing as the GPU but on the CPU, even if significant slower
- Using multiple processes using the same GPU needs additional work
 - 1 process = 1 context on the GPU – Contexts cannot run in parallel on the GPU
 - Multi-process Service (MPS) will time multiplex all calls of N processes into one context



CUDA Streams – Some more notes ++

- Additional knowledge – not needed for exam ;-)
- Use modern hardware / CUDA versions / sm-arch or things will get more complicated ...



* source: <https://developer.download.nvidia.com/CUDA/training/StreamsAndConcurrencyWebinar.pdf> nvidia.com